

=== CIFAR-10 Classification: Optimizer and Scheduler Comparison ===  
Using device: cuda

Original







Random Crop



Horizontal Flip



Rotat



=== Running Experiment: SGD\_NoScheduler ===

Model parameters: 1,149,770

Using SGD optimizer with learning rate: 0.01

No learning rate scheduler

[Epoch: 1/50] Loss: 1.5609, Accuracy: 42.46%, Learning rate: 0.010000,  
Time: 24.15s

[Epoch: 2/50] Loss: 1.1763, Accuracy: 58.06%, Learning rate: 0.010000,  
Time: 23.82s

[Epoch: 3/50] Loss: 1.0097, Accuracy: 64.23%, Learning rate: 0.010000,  
Time: 24.00s

[Epoch: 4/50] Loss: 0.9044, Accuracy: 68.18%, Learning rate: 0.010000,  
Time: 23.73s

[Epoch: 5/50] Loss: 0.8305, Accuracy: 71.01%, Learning rate: 0.010000,  
Time: 24.32s

[Epoch: 6/50] Loss: 0.7727, Accuracy: 73.17%, Learning rate: 0.010000,  
Time: 23.89s

[Epoch: 7/50] Loss: 0.7275, Accuracy: 74.88%, Learning rate: 0.010000,  
Time: 23.74s

[Epoch: 8/50] Loss: 0.6923, Accuracy: 76.09%, Learning rate: 0.010000,  
Time: 24.69s

[Epoch: 9/50] Loss: 0.6587, Accuracy: 77.17%, Learning rate: 0.010000,  
Time: 24.03s

[Epoch: 10/50] Loss: 0.6287, Accuracy: 78.46%, Learning rate: 0.010000,  
Time: 24.23s

[Epoch: 11/50] Loss: 0.6094, Accuracy: 78.92%, Learning rate: 0.010000,  
Time: 24.12s

[Epoch: 12/50] Loss: 0.5829, Accuracy: 79.89%, Learning rate: 0.010000,  
Time: 23.71s

[Epoch: 13/50] Loss: 0.5676, Accuracy: 80.25%, Learning rate: 0.010000,  
Time: 23.82s

[Epoch: 14/50] Loss: 0.5523, Accuracy: 81.06%, Learning rate: 0.010000,  
Time: 23.98s

[Epoch: 15/50] Loss: 0.5323, Accuracy: 81.68%, Learning rate: 0.010000,  
Time: 23.74s

[Epoch: 16/50] Loss: 0.5192, Accuracy: 82.07%, Learning rate: 0.010000,  
Time: 23.87s

[Epoch: 17/50] Loss: 0.4962, Accuracy: 82.85%, Learning rate: 0.010000,  
Time: 23.75s

[Epoch: 18/50] Loss: 0.4852, Accuracy: 83.22%, Learning rate: 0.010000,  
Time: 24.03s

[Epoch: 19/50] Loss: 0.4727, Accuracy: 83.61%, Learning rate: 0.010000,  
Time: 23.75s

[Epoch: 20/50] Loss: 0.4664, Accuracy: 83.97%, Learning rate: 0.010000,  
Time: 23.78s  
[Epoch: 21/50] Loss: 0.4542, Accuracy: 84.42%, Learning rate: 0.010000,  
Time: 23.66s  
[Epoch: 22/50] Loss: 0.4457, Accuracy: 84.43%, Learning rate: 0.010000,  
Time: 23.75s  
[Epoch: 23/50] Loss: 0.4328, Accuracy: 85.05%, Learning rate: 0.010000,  
Time: 24.04s  
[Epoch: 24/50] Loss: 0.4296, Accuracy: 85.30%, Learning rate: 0.010000,  
Time: 23.86s  
[Epoch: 25/50] Loss: 0.4173, Accuracy: 85.52%, Learning rate: 0.010000,  
Time: 23.82s  
[Epoch: 26/50] Loss: 0.4098, Accuracy: 85.69%, Learning rate: 0.010000,  
Time: 23.84s  
[Epoch: 27/50] Loss: 0.4036, Accuracy: 86.19%, Learning rate: 0.010000,  
Time: 24.17s  
[Epoch: 28/50] Loss: 0.3982, Accuracy: 86.25%, Learning rate: 0.010000,  
Time: 23.90s  
[Epoch: 29/50] Loss: 0.3848, Accuracy: 86.82%, Learning rate: 0.010000,  
Time: 23.78s  
[Epoch: 30/50] Loss: 0.3823, Accuracy: 86.83%, Learning rate: 0.010000,  
Time: 23.95s  
[Epoch: 31/50] Loss: 0.3728, Accuracy: 87.20%, Learning rate: 0.010000,  
Time: 23.79s  
[Epoch: 32/50] Loss: 0.3681, Accuracy: 87.32%, Learning rate: 0.010000,  
Time: 24.38s  
[Epoch: 33/50] Loss: 0.3638, Accuracy: 87.39%, Learning rate: 0.010000,  
Time: 23.93s  
[Epoch: 34/50] Loss: 0.3589, Accuracy: 87.54%, Learning rate: 0.010000,  
Time: 24.31s  
[Epoch: 35/50] Loss: 0.3503, Accuracy: 87.96%, Learning rate: 0.010000,  
Time: 24.31s  
[Epoch: 36/50] Loss: 0.3476, Accuracy: 88.03%, Learning rate: 0.010000,  
Time: 24.21s  
[Epoch: 37/50] Loss: 0.3422, Accuracy: 88.19%, Learning rate: 0.010000,  
Time: 23.75s  
[Epoch: 38/50] Loss: 0.3400, Accuracy: 88.23%, Learning rate: 0.010000,  
Time: 23.66s  
[Epoch: 39/50] Loss: 0.3375, Accuracy: 88.31%, Learning rate: 0.010000,  
Time: 24.08s  
[Epoch: 40/50] Loss: 0.3291, Accuracy: 88.62%, Learning rate: 0.010000,  
Time: 23.75s

[Epoch: 41/50] Loss: 0.3237, Accuracy: 88.94%, Learning rate: 0.010000,  
Time: 23.81s  
[Epoch: 42/50] Loss: 0.3243, Accuracy: 88.82%, Learning rate: 0.010000,  
Time: 23.93s  
[Epoch: 43/50] Loss: 0.3234, Accuracy: 88.72%, Learning rate: 0.010000,  
Time: 24.04s  
[Epoch: 44/50] Loss: 0.3194, Accuracy: 88.81%, Learning rate: 0.010000,  
Time: 23.76s  
[Epoch: 45/50] Loss: 0.3125, Accuracy: 89.20%, Learning rate: 0.010000,  
Time: 23.80s  
[Epoch: 46/50] Loss: 0.3062, Accuracy: 89.51%, Learning rate: 0.010000,  
Time: 23.77s  
[Epoch: 47/50] Loss: 0.3090, Accuracy: 89.38%, Learning rate: 0.010000,  
Time: 23.92s  
[Epoch: 48/50] Loss: 0.3125, Accuracy: 89.14%, Learning rate: 0.010000,  
Time: 24.01s  
[Epoch: 49/50] Loss: 0.2990, Accuracy: 89.50%, Learning rate: 0.010000,  
Time: 23.75s  
[Epoch: 50/50] Loss: 0.2998, Accuracy: 89.61%, Learning rate: 0.010000,  
Time: 23.80s  
Test set accuracy: 86.92%

=== Running Experiment: SGD\_StepLR ===

Model parameters: 1,149,770

Using SGD optimizer with learning rate: 0.01

Using StepLR scheduler with step\_size=15, gamma=0.1

[Epoch: 1/50] Loss: 1.5582, Accuracy: 42.62%, Learning rate: 0.010000,  
Time: 24.06s  
[Epoch: 2/50] Loss: 1.1810, Accuracy: 57.44%, Learning rate: 0.010000,  
Time: 24.11s  
[Epoch: 3/50] Loss: 1.0199, Accuracy: 63.83%, Learning rate: 0.010000,  
Time: 23.95s  
[Epoch: 4/50] Loss: 0.9114, Accuracy: 67.87%, Learning rate: 0.010000,  
Time: 24.02s  
[Epoch: 5/50] Loss: 0.8354, Accuracy: 71.05%, Learning rate: 0.010000,  
Time: 24.36s  
[Epoch: 6/50] Loss: 0.7748, Accuracy: 73.14%, Learning rate: 0.010000,  
Time: 23.91s  
[Epoch: 7/50] Loss: 0.7291, Accuracy: 74.51%, Learning rate: 0.010000,  
Time: 24.03s  
[Epoch: 8/50] Loss: 0.6879, Accuracy: 76.06%, Learning rate: 0.010000,  
Time: 23.93s

[Epoch: 9/50] Loss: 0.6584, Accuracy: 77.39%, Learning rate: 0.010000,  
Time: 24.08s  
[Epoch: 10/50] Loss: 0.6339, Accuracy: 78.18%, Learning rate: 0.010000,  
Time: 24.23s  
[Epoch: 11/50] Loss: 0.6031, Accuracy: 79.27%, Learning rate: 0.010000,  
Time: 24.03s  
[Epoch: 12/50] Loss: 0.5797, Accuracy: 79.78%, Learning rate: 0.010000,  
Time: 23.97s  
[Epoch: 13/50] Loss: 0.5604, Accuracy: 80.52%, Learning rate: 0.010000,  
Time: 23.89s  
[Epoch: 14/50] Loss: 0.5470, Accuracy: 81.14%, Learning rate: 0.010000,  
Time: 23.92s  
[Epoch: 15/50] Loss: 0.5284, Accuracy: 81.69%, Learning rate: 0.001000,  
Time: 23.84s  
[Epoch: 16/50] Loss: 0.4419, Accuracy: 85.03%, Learning rate: 0.001000,  
Time: 23.73s  
[Epoch: 17/50] Loss: 0.4243, Accuracy: 85.55%, Learning rate: 0.001000,  
Time: 24.52s  
[Epoch: 18/50] Loss: 0.4108, Accuracy: 85.81%, Learning rate: 0.001000,  
Time: 24.11s  
[Epoch: 19/50] Loss: 0.4055, Accuracy: 86.07%, Learning rate: 0.001000,  
Time: 24.31s  
[Epoch: 20/50] Loss: 0.3990, Accuracy: 86.38%, Learning rate: 0.001000,  
Time: 24.48s  
[Epoch: 21/50] Loss: 0.3915, Accuracy: 86.66%, Learning rate: 0.001000,  
Time: 24.23s  
[Epoch: 22/50] Loss: 0.3894, Accuracy: 86.75%, Learning rate: 0.001000,  
Time: 24.06s  
[Epoch: 23/50] Loss: 0.3869, Accuracy: 86.66%, Learning rate: 0.001000,  
Time: 24.12s  
[Epoch: 24/50] Loss: 0.3813, Accuracy: 86.89%, Learning rate: 0.001000,  
Time: 23.86s  
[Epoch: 25/50] Loss: 0.3813, Accuracy: 87.00%, Learning rate: 0.001000,  
Time: 23.72s  
[Epoch: 26/50] Loss: 0.3760, Accuracy: 87.08%, Learning rate: 0.001000,  
Time: 23.93s  
[Epoch: 27/50] Loss: 0.3766, Accuracy: 87.07%, Learning rate: 0.001000,  
Time: 23.97s  
[Epoch: 28/50] Loss: 0.3689, Accuracy: 87.44%, Learning rate: 0.001000,  
Time: 23.83s  
[Epoch: 29/50] Loss: 0.3694, Accuracy: 87.37%, Learning rate: 0.001000,  
Time: 24.36s

[Epoch: 30/50] Loss: 0.3665, Accuracy: 87.30%, Learning rate: 0.000100,  
Time: 24.05s  
[Epoch: 31/50] Loss: 0.3519, Accuracy: 87.96%, Learning rate: 0.000100,  
Time: 23.90s  
[Epoch: 32/50] Loss: 0.3494, Accuracy: 87.95%, Learning rate: 0.000100,  
Time: 24.27s  
[Epoch: 33/50] Loss: 0.3487, Accuracy: 88.10%, Learning rate: 0.000100,  
Time: 24.49s  
[Epoch: 34/50] Loss: 0.3487, Accuracy: 88.17%, Learning rate: 0.000100,  
Time: 24.47s  
[Epoch: 35/50] Loss: 0.3455, Accuracy: 88.17%, Learning rate: 0.000100,  
Time: 24.11s  
[Epoch: 36/50] Loss: 0.3433, Accuracy: 88.22%, Learning rate: 0.000100,  
Time: 23.60s  
[Epoch: 37/50] Loss: 0.3448, Accuracy: 88.26%, Learning rate: 0.000100,  
Time: 24.17s  
[Epoch: 38/50] Loss: 0.3437, Accuracy: 88.14%, Learning rate: 0.000100,  
Time: 23.85s  
[Epoch: 39/50] Loss: 0.3444, Accuracy: 88.20%, Learning rate: 0.000100,  
Time: 23.76s  
[Epoch: 40/50] Loss: 0.3432, Accuracy: 88.16%, Learning rate: 0.000100,  
Time: 24.11s  
[Epoch: 41/50] Loss: 0.3422, Accuracy: 88.24%, Learning rate: 0.000100,  
Time: 23.81s  
[Epoch: 42/50] Loss: 0.3398, Accuracy: 88.21%, Learning rate: 0.000100,  
Time: 23.83s  
[Epoch: 43/50] Loss: 0.3442, Accuracy: 88.21%, Learning rate: 0.000100,  
Time: 23.70s  
[Epoch: 44/50] Loss: 0.3420, Accuracy: 88.31%, Learning rate: 0.000100,  
Time: 23.77s  
[Epoch: 45/50] Loss: 0.3410, Accuracy: 88.24%, Learning rate: 0.000010,  
Time: 24.26s  
[Epoch: 46/50] Loss: 0.3365, Accuracy: 88.40%, Learning rate: 0.000010,  
Time: 23.96s  
[Epoch: 47/50] Loss: 0.3431, Accuracy: 88.34%, Learning rate: 0.000010,  
Time: 24.00s  
[Epoch: 48/50] Loss: 0.3349, Accuracy: 88.56%, Learning rate: 0.000010,  
Time: 23.97s  
[Epoch: 49/50] Loss: 0.3379, Accuracy: 88.34%, Learning rate: 0.000010,  
Time: 23.81s  
[Epoch: 50/50] Loss: 0.3409, Accuracy: 88.37%, Learning rate: 0.000010,  
Time: 24.08s  
Test set accuracy: 88.57%

=== Running Experiment: SGD\_CosineAnnealing ===

Model parameters: 1,149,770

Using SGD optimizer with learning rate: 0.01

Using CosineAnnealingLR scheduler with T\_max=50

[Epoch: 1/50] Loss: 1.5585, Accuracy: 42.59%, Learning rate: 0.009990,  
Time: 24.47s

[Epoch: 2/50] Loss: 1.1820, Accuracy: 57.49%, Learning rate: 0.009961,  
Time: 24.09s

[Epoch: 3/50] Loss: 1.0252, Accuracy: 63.70%, Learning rate: 0.009911,  
Time: 23.79s

[Epoch: 4/50] Loss: 0.9264, Accuracy: 67.31%, Learning rate: 0.009843,  
Time: 23.78s

[Epoch: 5/50] Loss: 0.8398, Accuracy: 70.61%, Learning rate: 0.009755,  
Time: 24.33s

[Epoch: 6/50] Loss: 0.7797, Accuracy: 72.84%, Learning rate: 0.009649,  
Time: 24.16s

[Epoch: 7/50] Loss: 0.7331, Accuracy: 74.39%, Learning rate: 0.009524,  
Time: 24.20s

[Epoch: 8/50] Loss: 0.6903, Accuracy: 76.18%, Learning rate: 0.009382,  
Time: 23.91s

[Epoch: 9/50] Loss: 0.6656, Accuracy: 77.05%, Learning rate: 0.009222,  
Time: 24.15s

[Epoch: 10/50] Loss: 0.6367, Accuracy: 77.70%, Learning rate: 0.009045,  
Time: 23.83s

[Epoch: 11/50] Loss: 0.6058, Accuracy: 79.10%, Learning rate: 0.008853,  
Time: 23.93s

[Epoch: 12/50] Loss: 0.5797, Accuracy: 80.11%, Learning rate: 0.008645,  
Time: 23.72s

[Epoch: 13/50] Loss: 0.5648, Accuracy: 80.47%, Learning rate: 0.008423,  
Time: 23.86s

[Epoch: 14/50] Loss: 0.5400, Accuracy: 81.37%, Learning rate: 0.008187,  
Time: 23.85s

[Epoch: 15/50] Loss: 0.5202, Accuracy: 81.95%, Learning rate: 0.007939,  
Time: 24.44s

[Epoch: 16/50] Loss: 0.5066, Accuracy: 82.55%, Learning rate: 0.007679,  
Time: 23.77s

[Epoch: 17/50] Loss: 0.4939, Accuracy: 82.90%, Learning rate: 0.007409,  
Time: 23.88s

[Epoch: 18/50] Loss: 0.4786, Accuracy: 83.61%, Learning rate: 0.007129,  
Time: 23.85s

[Epoch: 19/50] Loss: 0.4650, Accuracy: 83.90%, Learning rate: 0.006841,  
Time: 23.78s



[Epoch: 20/50] Loss: 0.4484, Accuracy: 84.45%, Learning rate: 0.006545,  
Time: 23.76s  
[Epoch: 21/50] Loss: 0.4319, Accuracy: 85.14%, Learning rate: 0.006243,  
Time: 23.71s  
[Epoch: 22/50] Loss: 0.4240, Accuracy: 85.35%, Learning rate: 0.005937,  
Time: 24.05s  
[Epoch: 23/50] Loss: 0.4138, Accuracy: 85.86%, Learning rate: 0.005627,  
Time: 23.91s  
[Epoch: 24/50] Loss: 0.3974, Accuracy: 86.29%, Learning rate: 0.005314,  
Time: 24.01s  
[Epoch: 25/50] Loss: 0.3913, Accuracy: 86.49%, Learning rate: 0.005000,  
Time: 23.75s  
[Epoch: 26/50] Loss: 0.3770, Accuracy: 86.93%, Learning rate: 0.004686,  
Time: 24.09s  
[Epoch: 27/50] Loss: 0.3705, Accuracy: 87.11%, Learning rate: 0.004373,  
Time: 24.05s  
[Epoch: 28/50] Loss: 0.3611, Accuracy: 87.54%, Learning rate: 0.004063,  
Time: 23.86s  
[Epoch: 29/50] Loss: 0.3466, Accuracy: 88.02%, Learning rate: 0.003757,  
Time: 23.92s  
[Epoch: 30/50] Loss: 0.3416, Accuracy: 88.25%, Learning rate: 0.003455,  
Time: 23.86s  
[Epoch: 31/50] Loss: 0.3297, Accuracy: 88.50%, Learning rate: 0.003159,  
Time: 24.47s  
[Epoch: 32/50] Loss: 0.3201, Accuracy: 88.88%, Learning rate: 0.002871,  
Time: 23.88s  
[Epoch: 33/50] Loss: 0.3111, Accuracy: 89.45%, Learning rate: 0.002591,  
Time: 23.78s  
[Epoch: 34/50] Loss: 0.3036, Accuracy: 89.60%, Learning rate: 0.002321,  
Time: 24.61s  
[Epoch: 35/50] Loss: 0.2951, Accuracy: 89.79%, Learning rate: 0.002061,  
Time: 23.81s  
[Epoch: 36/50] Loss: 0.2875, Accuracy: 90.26%, Learning rate: 0.001813,  
Time: 23.77s  
[Epoch: 37/50] Loss: 0.2770, Accuracy: 90.65%, Learning rate: 0.001577,  
Time: 23.90s  
[Epoch: 38/50] Loss: 0.2717, Accuracy: 90.63%, Learning rate: 0.001355,  
Time: 23.90s  
[Epoch: 39/50] Loss: 0.2662, Accuracy: 90.81%, Learning rate: 0.001147,  
Time: 24.19s  
[Epoch: 40/50] Loss: 0.2569, Accuracy: 91.20%, Learning rate: 0.000955,  
Time: 23.73s

[Epoch: 41/50] Loss: 0.2526, Accuracy: 91.31%, Learning rate: 0.000778,  
Time: 24.05s  
[Epoch: 42/50] Loss: 0.2487, Accuracy: 91.45%, Learning rate: 0.000618,  
Time: 23.74s  
[Epoch: 43/50] Loss: 0.2436, Accuracy: 91.61%, Learning rate: 0.000476,  
Time: 24.22s  
[Epoch: 44/50] Loss: 0.2355, Accuracy: 91.98%, Learning rate: 0.000351,  
Time: 24.14s  
[Epoch: 45/50] Loss: 0.2345, Accuracy: 92.03%, Learning rate: 0.000245,  
Time: 23.84s  
[Epoch: 46/50] Loss: 0.2313, Accuracy: 92.07%, Learning rate: 0.000157,  
Time: 23.72s  
[Epoch: 47/50] Loss: 0.2294, Accuracy: 92.15%, Learning rate: 0.000089,  
Time: 23.89s  
[Epoch: 48/50] Loss: 0.2257, Accuracy: 92.39%, Learning rate: 0.000039,  
Time: 23.76s  
[Epoch: 49/50] Loss: 0.2263, Accuracy: 92.22%, Learning rate: 0.000010,  
Time: 23.74s  
[Epoch: 50/50] Loss: 0.2236, Accuracy: 92.32%, Learning rate: 0.000000,  
Time: 24.34s  
Test set accuracy: 90.33%

=== Running Experiment: Adam\_NoScheduler ===

Model parameters: 1,149,770

Using Adam optimizer with learning rate: 0.001

No learning rate scheduler

[Epoch: 1/50] Loss: 1.4612, Accuracy: 46.53%, Learning rate: 0.001000,  
Time: 23.74s  
[Epoch: 2/50] Loss: 1.1000, Accuracy: 61.16%, Learning rate: 0.001000,  
Time: 24.09s  
[Epoch: 3/50] Loss: 0.9401, Accuracy: 67.00%, Learning rate: 0.001000,  
Time: 24.07s  
[Epoch: 4/50] Loss: 0.8352, Accuracy: 70.97%, Learning rate: 0.001000,  
Time: 24.09s  
[Epoch: 5/50] Loss: 0.7672, Accuracy: 73.27%, Learning rate: 0.001000,  
Time: 24.27s  
[Epoch: 6/50] Loss: 0.7132, Accuracy: 75.45%, Learning rate: 0.001000,  
Time: 24.30s  
[Epoch: 7/50] Loss: 0.6755, Accuracy: 76.55%, Learning rate: 0.001000,  
Time: 24.10s  
[Epoch: 8/50] Loss: 0.6403, Accuracy: 78.02%, Learning rate: 0.001000,  
Time: 23.77s

[Epoch: 9/50] Loss: 0.6119, Accuracy: 79.04%, Learning rate: 0.001000,  
Time: 23.89s  
[Epoch: 10/50] Loss: 0.5938, Accuracy: 79.66%, Learning rate: 0.001000,  
Time: 23.84s  
[Epoch: 11/50] Loss: 0.5674, Accuracy: 80.50%, Learning rate: 0.001000,  
Time: 24.18s  
[Epoch: 12/50] Loss: 0.5501, Accuracy: 81.10%, Learning rate: 0.001000,  
Time: 23.86s  
[Epoch: 13/50] Loss: 0.5413, Accuracy: 81.51%, Learning rate: 0.001000,  
Time: 23.61s  
[Epoch: 14/50] Loss: 0.5245, Accuracy: 82.02%, Learning rate: 0.001000,  
Time: 23.98s  
[Epoch: 15/50] Loss: 0.5094, Accuracy: 82.76%, Learning rate: 0.001000,  
Time: 24.26s  
[Epoch: 16/50] Loss: 0.4979, Accuracy: 83.03%, Learning rate: 0.001000,  
Time: 23.88s  
[Epoch: 17/50] Loss: 0.4898, Accuracy: 83.41%, Learning rate: 0.001000,  
Time: 23.80s  
[Epoch: 18/50] Loss: 0.4836, Accuracy: 83.44%, Learning rate: 0.001000,  
Time: 24.53s  
[Epoch: 19/50] Loss: 0.4749, Accuracy: 83.82%, Learning rate: 0.001000,  
Time: 24.21s  
[Epoch: 20/50] Loss: 0.4658, Accuracy: 84.04%, Learning rate: 0.001000,  
Time: 24.00s  
[Epoch: 21/50] Loss: 0.4651, Accuracy: 84.26%, Learning rate: 0.001000,  
Time: 23.96s  
[Epoch: 22/50] Loss: 0.4558, Accuracy: 84.55%, Learning rate: 0.001000,  
Time: 23.94s  
[Epoch: 23/50] Loss: 0.4476, Accuracy: 84.82%, Learning rate: 0.001000,  
Time: 23.78s  
[Epoch: 24/50] Loss: 0.4449, Accuracy: 84.65%, Learning rate: 0.001000,  
Time: 24.15s  
[Epoch: 25/50] Loss: 0.4379, Accuracy: 85.05%, Learning rate: 0.001000,  
Time: 23.80s  
[Epoch: 26/50] Loss: 0.4386, Accuracy: 85.16%, Learning rate: 0.001000,  
Time: 23.82s  
[Epoch: 27/50] Loss: 0.4328, Accuracy: 85.21%, Learning rate: 0.001000,  
Time: 24.20s  
[Epoch: 28/50] Loss: 0.4234, Accuracy: 85.55%, Learning rate: 0.001000,  
Time: 23.89s  
[Epoch: 29/50] Loss: 0.4282, Accuracy: 85.31%, Learning rate: 0.001000,  
Time: 24.41s

[Epoch: 30/50] Loss: 0.4199, Accuracy: 85.70%, Learning rate: 0.001000,  
Time: 24.08s  
[Epoch: 31/50] Loss: 0.4141, Accuracy: 85.96%, Learning rate: 0.001000,  
Time: 23.82s  
[Epoch: 32/50] Loss: 0.4144, Accuracy: 85.83%, Learning rate: 0.001000,  
Time: 23.74s  
[Epoch: 33/50] Loss: 0.4145, Accuracy: 85.82%, Learning rate: 0.001000,  
Time: 23.72s  
[Epoch: 34/50] Loss: 0.4078, Accuracy: 86.19%, Learning rate: 0.001000,  
Time: 24.00s  
[Epoch: 35/50] Loss: 0.4074, Accuracy: 86.11%, Learning rate: 0.001000,  
Time: 23.85s  
[Epoch: 36/50] Loss: 0.4055, Accuracy: 86.24%, Learning rate: 0.001000,  
Time: 24.47s  
[Epoch: 37/50] Loss: 0.4012, Accuracy: 86.19%, Learning rate: 0.001000,  
Time: 23.97s  
[Epoch: 38/50] Loss: 0.3995, Accuracy: 86.38%, Learning rate: 0.001000,  
Time: 24.10s  
[Epoch: 39/50] Loss: 0.3961, Accuracy: 86.48%, Learning rate: 0.001000,  
Time: 23.83s  
[Epoch: 40/50] Loss: 0.3963, Accuracy: 86.73%, Learning rate: 0.001000,  
Time: 23.93s  
[Epoch: 41/50] Loss: 0.3921, Accuracy: 86.61%, Learning rate: 0.001000,  
Time: 23.69s  
[Epoch: 42/50] Loss: 0.3914, Accuracy: 86.56%, Learning rate: 0.001000,  
Time: 23.76s  
[Epoch: 43/50] Loss: 0.3918, Accuracy: 86.98%, Learning rate: 0.001000,  
Time: 23.70s  
[Epoch: 44/50] Loss: 0.3869, Accuracy: 86.74%, Learning rate: 0.001000,  
Time: 24.03s  
[Epoch: 45/50] Loss: 0.3876, Accuracy: 86.90%, Learning rate: 0.001000,  
Time: 24.38s  
[Epoch: 46/50] Loss: 0.3855, Accuracy: 86.91%, Learning rate: 0.001000,  
Time: 24.37s  
[Epoch: 47/50] Loss: 0.3851, Accuracy: 87.05%, Learning rate: 0.001000,  
Time: 23.89s  
[Epoch: 48/50] Loss: 0.3852, Accuracy: 86.87%, Learning rate: 0.001000,  
Time: 23.72s  
[Epoch: 49/50] Loss: 0.3778, Accuracy: 87.22%, Learning rate: 0.001000,  
Time: 23.74s  
[Epoch: 50/50] Loss: 0.3782, Accuracy: 87.32%, Learning rate: 0.001000,  
Time: 23.96s  
Test set accuracy: 87.25%

=== Running Experiment: Adam\_StepLR ===

Model parameters: 1,149,770

Using Adam optimizer with learning rate: 0.001

Using StepLR scheduler with step\_size=15, gamma=0.1

[Epoch: 1/50] Loss: 1.4760, Accuracy: 45.89%, Learning rate: 0.001000,  
Time: 23.83s

[Epoch: 2/50] Loss: 1.0977, Accuracy: 60.84%, Learning rate: 0.001000,  
Time: 23.97s

[Epoch: 3/50] Loss: 0.9327, Accuracy: 67.41%, Learning rate: 0.001000,  
Time: 24.07s

[Epoch: 4/50] Loss: 0.8310, Accuracy: 71.15%, Learning rate: 0.001000,  
Time: 23.85s

[Epoch: 5/50] Loss: 0.7598, Accuracy: 73.71%, Learning rate: 0.001000,  
Time: 24.24s

[Epoch: 6/50] Loss: 0.7124, Accuracy: 75.27%, Learning rate: 0.001000,  
Time: 23.77s

[Epoch: 7/50] Loss: 0.6721, Accuracy: 76.92%, Learning rate: 0.001000,  
Time: 24.44s

[Epoch: 8/50] Loss: 0.6393, Accuracy: 78.14%, Learning rate: 0.001000,  
Time: 24.27s

[Epoch: 9/50] Loss: 0.6106, Accuracy: 78.91%, Learning rate: 0.001000,  
Time: 23.78s

[Epoch: 10/50] Loss: 0.5898, Accuracy: 79.71%, Learning rate: 0.001000,  
Time: 23.86s

[Epoch: 11/50] Loss: 0.5703, Accuracy: 80.33%, Learning rate: 0.001000,  
Time: 23.83s

[Epoch: 12/50] Loss: 0.5560, Accuracy: 80.93%, Learning rate: 0.001000,  
Time: 23.96s

[Epoch: 13/50] Loss: 0.5405, Accuracy: 81.68%, Learning rate: 0.001000,  
Time: 23.83s

[Epoch: 14/50] Loss: 0.5242, Accuracy: 81.92%, Learning rate: 0.001000,  
Time: 24.42s

[Epoch: 15/50] Loss: 0.5113, Accuracy: 82.47%, Learning rate: 0.000100,  
Time: 23.77s

[Epoch: 16/50] Loss: 0.4055, Accuracy: 86.15%, Learning rate: 0.000100,  
Time: 23.84s

[Epoch: 17/50] Loss: 0.3688, Accuracy: 87.42%, Learning rate: 0.000100,  
Time: 23.89s

[Epoch: 18/50] Loss: 0.3507, Accuracy: 88.12%, Learning rate: 0.000100,  
Time: 23.92s

[Epoch: 19/50] Loss: 0.3392, Accuracy: 88.36%, Learning rate: 0.000100,  
Time: 24.33s

[Epoch: 20/50] Loss: 0.3337, Accuracy: 88.64%, Learning rate: 0.000100,  
Time: 24.09s  
[Epoch: 21/50] Loss: 0.3181, Accuracy: 89.15%, Learning rate: 0.000100,  
Time: 24.11s  
[Epoch: 22/50] Loss: 0.3156, Accuracy: 89.33%, Learning rate: 0.000100,  
Time: 23.70s  
[Epoch: 23/50] Loss: 0.3093, Accuracy: 89.53%, Learning rate: 0.000100,  
Time: 23.80s  
[Epoch: 24/50] Loss: 0.2991, Accuracy: 89.86%, Learning rate: 0.000100,  
Time: 23.64s  
[Epoch: 25/50] Loss: 0.2957, Accuracy: 89.80%, Learning rate: 0.000100,  
Time: 23.89s  
[Epoch: 26/50] Loss: 0.2910, Accuracy: 90.05%, Learning rate: 0.000100,  
Time: 23.66s  
[Epoch: 27/50] Loss: 0.2832, Accuracy: 90.30%, Learning rate: 0.000100,  
Time: 23.81s  
[Epoch: 28/50] Loss: 0.2819, Accuracy: 90.49%, Learning rate: 0.000100,  
Time: 23.80s  
[Epoch: 29/50] Loss: 0.2726, Accuracy: 90.64%, Learning rate: 0.000100,  
Time: 23.94s  
[Epoch: 30/50] Loss: 0.2681, Accuracy: 90.82%, Learning rate: 0.000010,  
Time: 24.46s  
[Epoch: 31/50] Loss: 0.2518, Accuracy: 91.52%, Learning rate: 0.000010,  
Time: 24.27s  
[Epoch: 32/50] Loss: 0.2476, Accuracy: 91.55%, Learning rate: 0.000010,  
Time: 23.72s  
[Epoch: 33/50] Loss: 0.2478, Accuracy: 91.61%, Learning rate: 0.000010,  
Time: 23.78s  
[Epoch: 34/50] Loss: 0.2454, Accuracy: 91.50%, Learning rate: 0.000010,  
Time: 23.91s  
[Epoch: 35/50] Loss: 0.2462, Accuracy: 91.64%, Learning rate: 0.000010,  
Time: 24.08s  
[Epoch: 36/50] Loss: 0.2417, Accuracy: 91.73%, Learning rate: 0.000010,  
Time: 23.73s  
[Epoch: 37/50] Loss: 0.2403, Accuracy: 91.79%, Learning rate: 0.000010,  
Time: 24.07s  
[Epoch: 38/50] Loss: 0.2381, Accuracy: 91.83%, Learning rate: 0.000010,  
Time: 23.93s  
[Epoch: 39/50] Loss: 0.2414, Accuracy: 91.74%, Learning rate: 0.000010,  
Time: 24.01s  
[Epoch: 40/50] Loss: 0.2394, Accuracy: 91.88%, Learning rate: 0.000010,  
Time: 23.74s

[Epoch: 41/50] Loss: 0.2364, Accuracy: 91.96%, Learning rate: 0.000010,  
Time: 24.17s  
[Epoch: 42/50] Loss: 0.2358, Accuracy: 92.01%, Learning rate: 0.000010,  
Time: 24.33s  
[Epoch: 43/50] Loss: 0.2353, Accuracy: 91.99%, Learning rate: 0.000010,  
Time: 23.81s  
[Epoch: 44/50] Loss: 0.2369, Accuracy: 91.96%, Learning rate: 0.000010,  
Time: 24.01s  
[Epoch: 45/50] Loss: 0.2357, Accuracy: 91.91%, Learning rate: 0.000001,  
Time: 24.17s  
[Epoch: 46/50] Loss: 0.2308, Accuracy: 92.11%, Learning rate: 0.000001,  
Time: 24.02s  
[Epoch: 47/50] Loss: 0.2333, Accuracy: 92.07%, Learning rate: 0.000001,  
Time: 24.03s  
[Epoch: 48/50] Loss: 0.2347, Accuracy: 92.04%, Learning rate: 0.000001,  
Time: 23.85s  
[Epoch: 49/50] Loss: 0.2288, Accuracy: 92.19%, Learning rate: 0.000001,  
Time: 23.66s  
[Epoch: 50/50] Loss: 0.2331, Accuracy: 92.04%, Learning rate: 0.000001,  
Time: 24.46s  
Test set accuracy: 90.48%

=== Running Experiment: Adam\_CosineAnnealing ===

Model parameters: 1,149,770

Using Adam optimizer with learning rate: 0.001

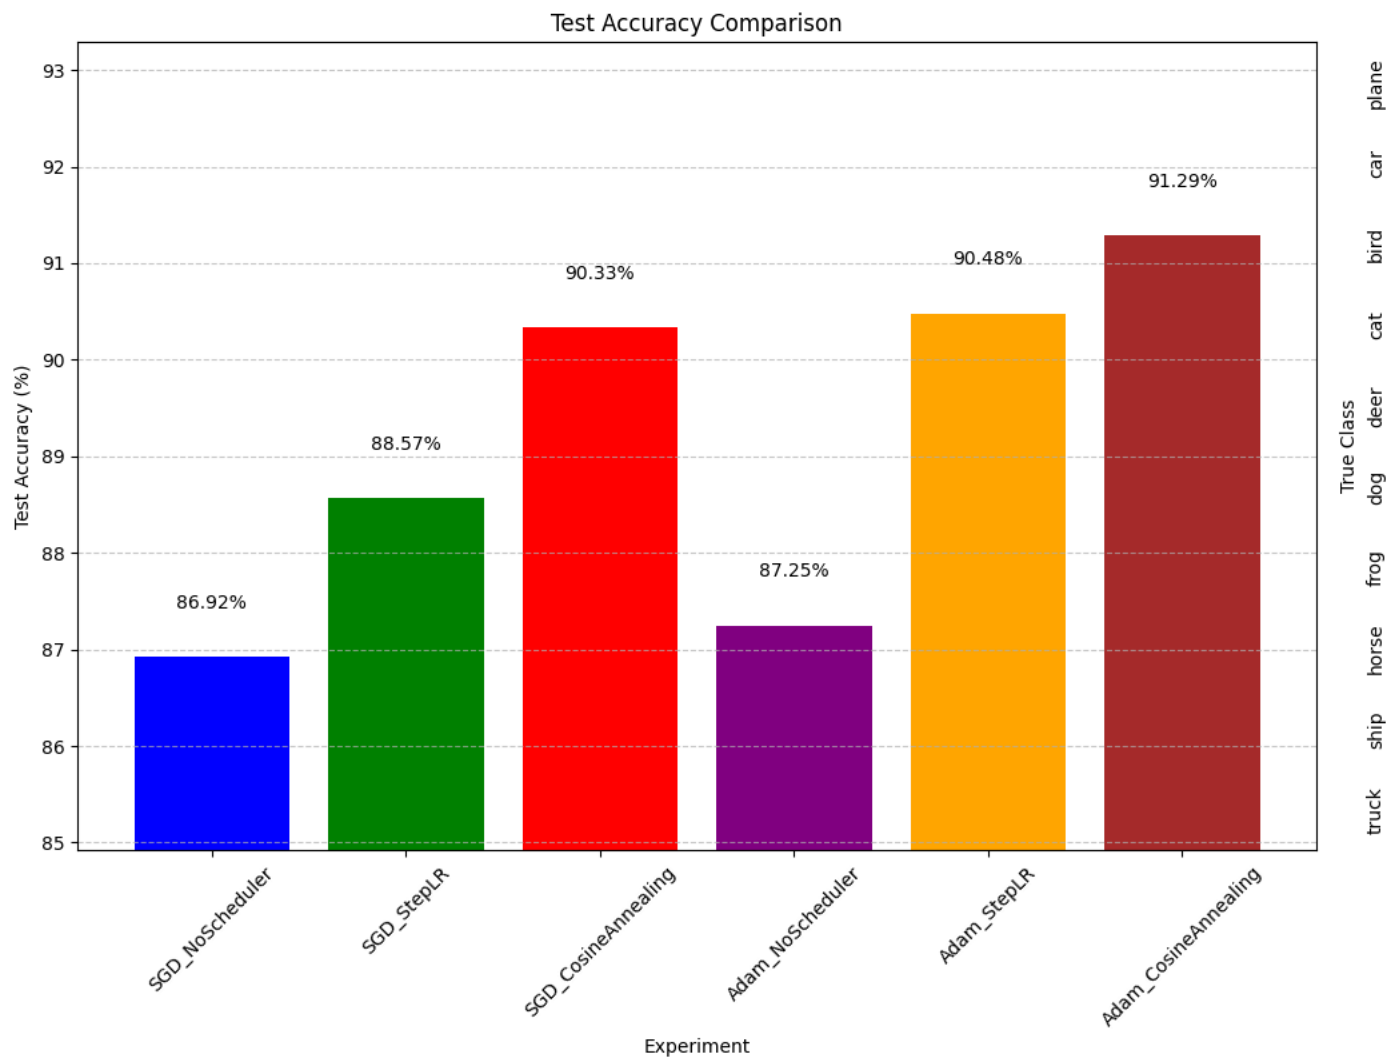

Using CosineAnnealingLR scheduler with T\_max=50

[Epoch: 1/50] Loss: 1.4764, Accuracy: 45.85%, Learning rate: 0.000999,  
Time: 24.30s  
[Epoch: 2/50] Loss: 1.0927, Accuracy: 61.50%, Learning rate: 0.000996,  
Time: 23.83s  
[Epoch: 3/50] Loss: 0.9319, Accuracy: 67.26%, Learning rate: 0.000991,  
Time: 23.75s  
[Epoch: 4/50] Loss: 0.8285, Accuracy: 71.06%, Learning rate: 0.000984,  
Time: 24.22s  
[Epoch: 5/50] Loss: 0.7588, Accuracy: 73.76%, Learning rate: 0.000976,  
Time: 24.31s  
[Epoch: 6/50] Loss: 0.7077, Accuracy: 75.66%, Learning rate: 0.000965,  
Time: 23.91s  
[Epoch: 7/50] Loss: 0.6696, Accuracy: 76.98%, Learning rate: 0.000952,  
Time: 23.73s  
[Epoch: 8/50] Loss: 0.6316, Accuracy: 78.44%, Learning rate: 0.000938,  
Time: 23.76s

[Epoch: 9/50] Loss: 0.6012, Accuracy: 79.47%, Learning rate: 0.000922,  
Time: 24.12s  
[Epoch: 10/50] Loss: 0.5750, Accuracy: 80.19%, Learning rate: 0.000905,  
Time: 23.89s  
[Epoch: 11/50] Loss: 0.5618, Accuracy: 80.73%, Learning rate: 0.000885,  
Time: 24.04s  
[Epoch: 12/50] Loss: 0.5376, Accuracy: 81.56%, Learning rate: 0.000864,  
Time: 24.57s  
[Epoch: 13/50] Loss: 0.5150, Accuracy: 82.28%, Learning rate: 0.000842,  
Time: 24.11s  
[Epoch: 14/50] Loss: 0.5034, Accuracy: 82.63%, Learning rate: 0.000819,  
Time: 24.23s  
[Epoch: 15/50] Loss: 0.4822, Accuracy: 83.37%, Learning rate: 0.000794,  
Time: 23.88s  
[Epoch: 16/50] Loss: 0.4736, Accuracy: 83.98%, Learning rate: 0.000768,  
Time: 23.79s  
[Epoch: 17/50] Loss: 0.4558, Accuracy: 84.37%, Learning rate: 0.000741,  
Time: 23.78s  
[Epoch: 18/50] Loss: 0.4430, Accuracy: 84.91%, Learning rate: 0.000713,  
Time: 24.20s  
[Epoch: 19/50] Loss: 0.4333, Accuracy: 85.11%, Learning rate: 0.000684,  
Time: 24.38s  
[Epoch: 20/50] Loss: 0.4190, Accuracy: 85.76%, Learning rate: 0.000655,  
Time: 24.01s  
[Epoch: 21/50] Loss: 0.4051, Accuracy: 86.20%, Learning rate: 0.000624,  
Time: 24.26s  
[Epoch: 22/50] Loss: 0.3944, Accuracy: 86.63%, Learning rate: 0.000594,  
Time: 24.35s  
[Epoch: 23/50] Loss: 0.3838, Accuracy: 86.91%, Learning rate: 0.000563,  
Time: 24.65s  
[Epoch: 24/50] Loss: 0.3697, Accuracy: 87.38%, Learning rate: 0.000531,  
Time: 24.19s  
[Epoch: 25/50] Loss: 0.3590, Accuracy: 87.82%, Learning rate: 0.000500,  
Time: 23.96s  
[Epoch: 26/50] Loss: 0.3486, Accuracy: 87.93%, Learning rate: 0.000469,  
Time: 23.94s  
[Epoch: 27/50] Loss: 0.3405, Accuracy: 88.41%, Learning rate: 0.000437,  
Time: 24.59s  
[Epoch: 28/50] Loss: 0.3255, Accuracy: 88.84%, Learning rate: 0.000406,  
Time: 24.66s  
[Epoch: 29/50] Loss: 0.3131, Accuracy: 89.10%, Learning rate: 0.000376,  
Time: 24.45s




[Epoch: 30/50] Loss: 0.3012, Accuracy: 89.69%, Learning rate: 0.000345,  
Time: 24.30s  
[Epoch: 31/50] Loss: 0.2925, Accuracy: 89.99%, Learning rate: 0.000316,  
Time: 24.27s  
[Epoch: 32/50] Loss: 0.2802, Accuracy: 90.25%, Learning rate: 0.000287,  
Time: 24.35s  
[Epoch: 33/50] Loss: 0.2724, Accuracy: 90.66%, Learning rate: 0.000259,  
Time: 24.33s  
[Epoch: 34/50] Loss: 0.2562, Accuracy: 91.17%, Learning rate: 0.000232,  
Time: 24.39s  
[Epoch: 35/50] Loss: 0.2442, Accuracy: 91.63%, Learning rate: 0.000206,  
Time: 24.80s  
[Epoch: 36/50] Loss: 0.2364, Accuracy: 91.92%, Learning rate: 0.000181,  
Time: 24.54s  
[Epoch: 37/50] Loss: 0.2230, Accuracy: 92.29%, Learning rate: 0.000158,  
Time: 24.51s  
[Epoch: 38/50] Loss: 0.2156, Accuracy: 92.58%, Learning rate: 0.000136,  
Time: 24.48s  
[Epoch: 39/50] Loss: 0.2076, Accuracy: 92.93%, Learning rate: 0.000115,  
Time: 24.38s  
[Epoch: 40/50] Loss: 0.2006, Accuracy: 93.18%, Learning rate: 0.000095,  
Time: 24.76s  
[Epoch: 41/50] Loss: 0.1932, Accuracy: 93.36%, Learning rate: 0.000078,  
Time: 24.74s  
[Epoch: 42/50] Loss: 0.1835, Accuracy: 93.84%, Learning rate: 0.000062,  
Time: 24.32s  
[Epoch: 43/50] Loss: 0.1783, Accuracy: 93.97%, Learning rate: 0.000048,  
Time: 24.32s  
[Epoch: 44/50] Loss: 0.1741, Accuracy: 94.09%, Learning rate: 0.000035,  
Time: 24.55s  
[Epoch: 45/50] Loss: 0.1693, Accuracy: 94.24%, Learning rate: 0.000024,  
Time: 24.24s  
[Epoch: 46/50] Loss: 0.1663, Accuracy: 94.34%, Learning rate: 0.000016,  
Time: 23.92s  
[Epoch: 47/50] Loss: 0.1624, Accuracy: 94.43%, Learning rate: 0.000009,  
Time: 24.67s  
[Epoch: 48/50] Loss: 0.1609, Accuracy: 94.71%, Learning rate: 0.000004,  
Time: 24.31s  
[Epoch: 49/50] Loss: 0.1573, Accuracy: 94.69%, Learning rate: 0.000001,  
Time: 24.86s  
[Epoch: 50/50] Loss: 0.1575, Accuracy: 94.69%, Learning rate: 0.000000,  
Time: 24.13s  
Test set accuracy: 91.29%









=== All Experiments Completed ===

SGD\_NoScheduler: Test Accuracy = 86.92%

SGD\_StepLR: Test Accuracy = 88.57%

SGD\_CosineAnnealing: Test Accuracy = 90.33%

Adam\_NoScheduler: Test Accuracy = 87.25%

Adam\_StepLR: Test Accuracy = 90.48%

Adam\_CosineAnnealing: Test Accuracy = 91.29%